

Deadlock Avoidance

Requires that the system has some additional information about how the resources are to be requested

- ☐ Simplest and most useful model requires that each process declare the *maximum number* of resources of each type that it may need.
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition.
- ☐ Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes.





Safe State

- ☐ When a process requests an available resource, system must decide if immediate allocation of those resources leaves the system in a *safe state*.
- System is in safe state if there exists a safe sequence of all processes.
- Sequence $\langle P_1, P_2, ..., P_n \rangle$ is safe if for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_i
 - If P_i resource needs are not immediately available, then P_i can wait until all P_i have finished.
 - When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate.
 - When P_i terminates, P_{i+1} can obtain its needed resources, and so on.





Basic Facts

- □ If a system is in safe state \Rightarrow no deadlocks.
- \square If a system is in unsafe state \Rightarrow possibility of deadlock.
- □ Avoidance ⇒ ensure that a system will never enter an unsafe state.





Safe State

□ System with 12 tape drives

•		Maximum Needs	Current Needs
	P0	10	5
	P1	4	2
	P2	9	2

At time t0: Safe sequence...<P1,P0,P2>

At time t1: Unsafe sequence if 1 more tape drive is allocated to P2



- Example 1: 12 instances of a resource
- □ At time t0, P0 holds 5, P1 holds 2, P2 holds 2
- \Box Available = 3 free instances

Example 1:

- 12 instances of a resource
- At time t0, P0 holds 5, P1 holds 2, P2 holds 2
- Available = 3 free instances

processes	max needs	allocated
P0	10	5
P1	4	2
P2	9	2



Example 1 (cont)

- Is the system in a safe state? Can I find a safe sequence?
- Yes, I claim the sequence <P1, P0, P2> is safe.
 - P1 requests its maximum (currently has 2, so needs 2 more) and holds 4, then there is only 1 free resource
 - Then P1 then releases all of its held resources, so that there are 5 free
 - Next, suppose P0 requests its maximum (currently has 5, so needs 5 more) and holds 10, so that there are 0 free
 - Then P0 releases all of its held resources, so that there are 10 free
 - Next P2 requests its max of 9, leaving 3 free and then releases them all
- Thus the sequence <P1, P0, P2> is safe, and is able in the worst-case to request maximum resources for each process in the sequence, and release all such resources for the next process in the sequence





Example 2:

at time t1, process P2 requests and is given 1
 more instance of the resource, then

processes	max needs	allocated
P0	10	5
P1	4	2
P2	9	3

- Available = 2 free instances
- Is the system in a safe state?

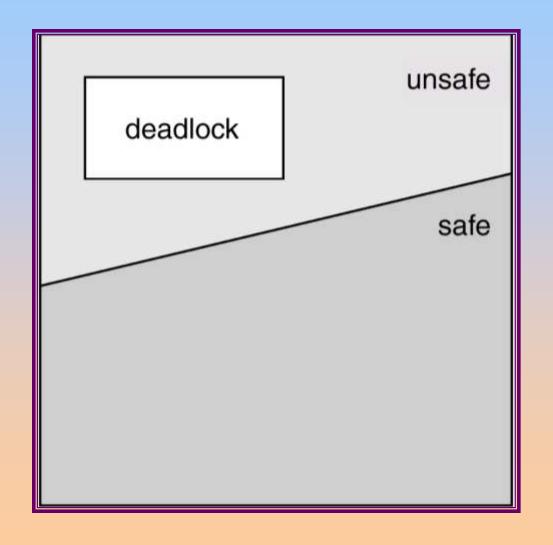


Example 2 (cont)

- P1 can request its maximum (currently holds 2, and needs 2 more) of 4, so that there are 0 free
- P1 releases all its held resources, so that available = 4 free
- Neither P0 nor P2 can request their maximum resources (P0 needs 5, P2 needs 6, and there are only 4 free)
 - Both would have to wait, so there could be deadlock
- The system is deemed unsafe
 - The mistake was granting P2 an extra resource at time t1
 - Forcing P2 to wait for P0 or P1 to release their resources would have avoided potential deadlock



Safe, Unsafe, Deadlock State

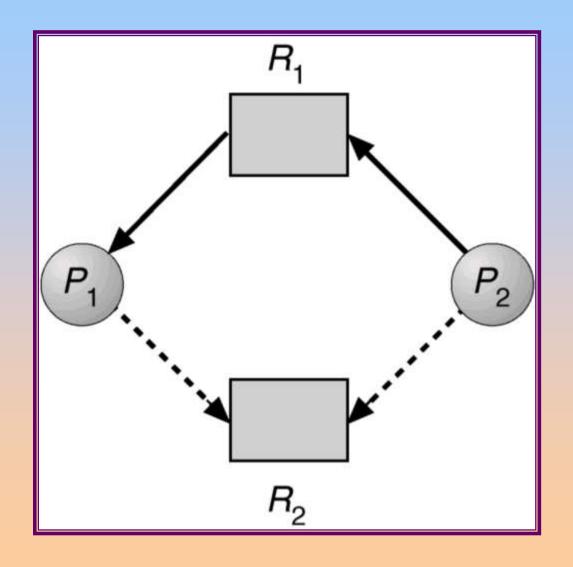




Resource-Allocation Graph Algorithm

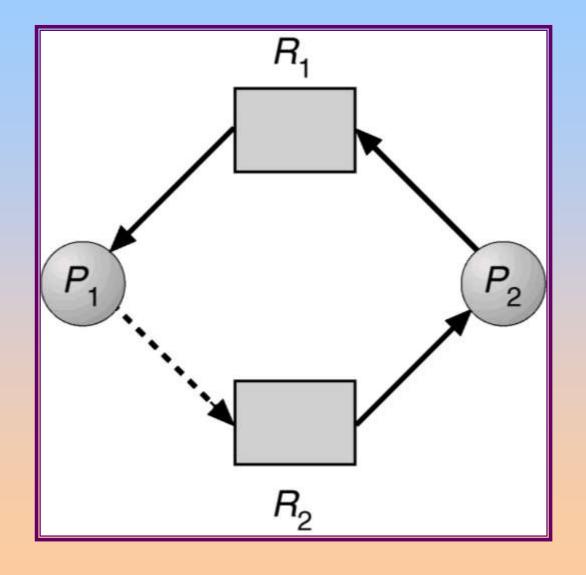
- SO 2: Demonstrate the resource allocation graph algorithm
- □ Claim edge $P_i \rightarrow R_j$ indicated that process P_i may request resource R_i ; represented by a dashed line.
- ☐ Claim edge converts to request edge when a process requests a resource.
- ☐ When a resource is released by a process, assignment edge reconverts to a claim edge.
- ☐ Resources must be claimed with a priori information in the system.

Resource-Allocation Graph For Deadlock Avoidance





Unsafe State In Resource-Allocation Graph





Banker's Algorithm

- Multiple instances.
- □ When a process enters into a system it should declare maximum resource needs in all resource types
- ☐ This number should not exceed total number of resources in the system.
- When a process requests a set of resources, the system must determine whether the allocation of these resources will leave the system in a safe state.
- ☐ If it will, the resources are allocated, otherwise it has to wait until some other process releases enough resources



Data Structures for the Banker's Algorithm

Let n = number of processes, and m = number of resources types.

- □ Available: Vector of length m. If available [j] = k, there are k instances of resource type R_i available.
- □ $Max: n \times m$ matrix. If Max[i,j] = k, then process P_i may request at most k instances of resource type R_{i} .
- □ Allocation: $n \times m$ matrix. If Allocation[i,j] = k then P_i is currently allocated k instances of $R_{j.}$
- □ Need: $n \times m$ matrix. If Need[i,j] = k, then P_i may need k more instances of R_i to complete its task.

Need[i,j] = Max[i,j] - Allocation[i,j].





Safety Algorithm

1. Let *Work* and *Finish* be vectors of length *m* and *n*, respectively. Initialize:

Work = Available
Finish
$$[i]$$
 = false for $i = 1, 2, 3, ..., n$.

- 2. Find an *i* such that both:
 - (a) Finish[i] = false
 - (b) *Need_i* ≤ *Work*

If no such *i* exists, go to step 4.

- 3. $Work = Work + Allocation_i$ Finish[i] = truego to step 2.
- 4. If *Finish* [*i*] == true for all *i*, then the system is in a safe state.

Resource-Request Algorithm for Process P_i

Let $Request_i$ be the request vector for process P_i . If $Request_i[j] = k$ then process P_i wants k instances of resource type R_i .

- 1. If *Request_i* ≤ *Need_i* go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
- 2. If $Request_i \le Available$, go to step 3. Otherwise P_i must wait, since resources are not available.
- 3. Pretend to allocate requested resources to P_i by modifying the state as follows:

```
Available = Available - Request<sub>i</sub>;

Allocation<sub>i</sub> = Allocation<sub>i</sub> + Request<sub>i</sub>;

Need<sub>i</sub> = Need<sub>i</sub> - Request<sub>i</sub>;
```

- If safe ⇒ the resources are allocated to P_i.
- If unsafe ⇒ P_i must wait, and the old resource-allocation state is restored





Example of Banker's Algorithm

- □ 5 processes P₀ through P₄; 3 resource types A (10 instances),
 B (5instances, and C (7 instances).
- \square Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	ABC	ABC	ABC
P_0	010	753	3 3 2
P_1	200	322	
P_2	302	902	
P_3	211	222	
P_4	002	433	



Example (Cont.)

The content of the matrix. Need is defined to be Max – Allocation.

$$\frac{Need}{ABC}$$
 P_0 743
 P_1 122
 P_2 600
 P_3 011
 P_4 431

☐ The system is in a safe state since the sequence $< P_1, P_3, P_4, P_2, P_0>$ satisfies safety criteria.

Solution: Banker's Algorithm

Step 1:

- □ Safety for process P₀
- \square need₀ = (7, 4, 3)
- □ If $need_0 \le Available$
- \square if $[(7, 4, 3) \le (3, 3, 2)]$ (false)
- ☐ Process P₀ must wait.

Step 2:

- □ Safety for process P₁
- \square need₁ = (1, 2, 2)
- \square if $need_i \leq Available$
- \Box if [(1, 2, 2) \le (3, 3, 2)]
- ☐ P_i will execute.
- □ Available = Available + Allocation
- \Box = (3, 3, 2) + (2, 0, 0)
- $\Box = (5, 3, 2)$

<u>Need</u>

- ABC
- $P_0 = 743$
- P_1 122
- $P_2 = 600$
- P₃ 011
- P₄ 431

Step 3:

- Safety for process P₂
- \Box need₂ = (6, 0, 0)
- ☐ if need₂ ≤ Available
- \square if [(6, 0, 0) \leq (5, 3, 2)] (false)
- \square P₃ will execute.
- □ Available = Available + Allocation
- \Box = (5, 3, 2) + (2, 1, 1)= (7, 4, 3)

Step 5:

- □ Safety for process P₄
- \Box need₄ = (4, 3, 1)
- ☐ If need₄ ≤ Available
- \Box If [(4, 3, 1) \leq (6, 4, 3)]
- □ P₄ will execute.
- □ Available = Available + Allocation
- $\Box = (7, 4, 3) + (0, 0, 2)$
- $\Box = (7, 4, 5)$

Step 6:

- Safety for process P_0 need₀ = (7, 4, 3)
- \square if $need_0 \leq Available$
- \Box if [(7, 4, 3) \leq (7, 4, 5)]
- \square P₀ will execute.
- □ Available = Available + Allocation
- \Box = (7, 4, 5) + (0, 1, 0)= (7, 5, 5)

Step 7:

- ☐ Safety for process P₂
- \square need₂ = (6, 0, 0)
- ☐ if need₂ ≤ Available
- \Box if [(6, 0, 0) \le (7, 5, 5)]
- \square P₂ will execute.
- □ Available = Available + Allocation
- \Box = (7, 5, 5) + (3, 0, 2)= (10, 5, 7)

Safety Sequence = <P1, P3, P4, P0, P2>



Example P_1 Request (1,0,2) (Cont.)

□ Check that Request \leq Available (that is, $(1,0,2) \leq (3,3,2) \Rightarrow true$.

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	ABC	ABC	ABC
P_0	010	7 4 3	230
P_1	302	020	
P_2	3 0 1	600	
P_3	211	011	
P_4	002	431	

□ Executing safety algorithm shows that sequence $< P_1, P_3, P_4, P_0, P_2>$ satisfies safety requirement.





Problems

Consider the following snapshot of a system.

	Allocation	Max	Available
	ABCD	ABCD	ABCD
P0	0012	0012	1520
P1	1000	1750	
P2	1354	2356	
P3	0632	0652	
P4	0014	0656	

Answer the following questions using the Banker's algorithm:

- a. What is the content of the matrix *Need?*
- b. Is the system in a safe state?
- c. If a request from process P1 arrives for (0,4,2,0), can the request be granted immediately?