

Illustrate the working principles of semaphore in critical section

Critical-Section Handling in OS

- Two approaches depending on if kernel is pre-emptive or non pre-emptive
 - Pre-emptive – allows pre-emption of process when running in kernel mode
 - Non-preemptive – runs until exits kernel mode, blocks, or voluntarily yields CPU
- Essentially free of race conditions in kernel mode

Synchronization Hardware

- Many systems provide hardware support for implementing the critical section code.
- All solutions below based on idea of locking

Protecting critical regions via locks

- Uniprocessors – could disable interrupts

Currently running code would execute without
preemption

Generally too inefficient on multiprocessor systems

Operating systems using this not broadly scalable

Synchronization Hardware

- Modern machines provide special atomic hardware instructions
- Atomic = non-interruptible
- Either test memory word and set value Or swap contents of two memory words

Solution to Critical-section Problem Using Locks

```
do  
{  
  acquire lock  
  critical section  
  release lock  
  remainder section  
} while (TRUE);
```

test_and_set Instruction

Definition:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

1. Executed atomically
2. Returns the original value of passed parameter
3. Set the new value of passed parameter to "TRUE".

Solution using test_and_set()

- Shared Boolean variable lock, initialized to FALSE
- Solution:

```
do {  
    while (test_and_set(&lock));  
        /* do nothing */  
        /* critical section */  
        lock = false;  
        /* remainder section */  
} while (true);
```

compare_and_swap Instruction

```
void swap(boolean *a,  
boolean *b)  
{  
    boolean temp;  
    temp = *b;  
    *b = *a;  
    *a = temp;  
}
```

- Mutual Exclusion with swap function

```
do {  
    Key = true;  
    While (key==true)  
        Swap (&lock,&key);  
    //Critical section  
    Lock = false;  
}while = true;
```


Mutex Locks

- Lock State & Operations
- Locks have two states
 - Held** : Someone in the critical section
 - Not held**: Nobody in the critical section
- Two Operation :
 - Acquire** : Mark lock as held or wait until release
 - Release** : Mark lock as un held

Mutex Locks

- Previous solutions are complicated and generally inaccessible to application programmers.
- OS designers build software tools to solve critical section problem
- Simplest is mutex lock .
- Protect a critical section by first acquire() a lock then release() the lock
- Boolean variable indicating if lock is available or not
- Calls to acquire() and release() must be atomic
- Usually implemented via hardware atomic instructions

acquire() and release()

```
acquire() {  
    while (!available);  
    /* busy wait */  
    available = false;;  
}  
release()  
{  
    available = true;  
}  
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
}  
while (true);
```

Semaphore

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to synchronize their activities.
- Semaphore **S** – integer variable
- Can only be accessed via two indivisible (atomic) operations

- **wait()** and **signal()**

- ▶ Originally called **P()** and **V()**

- Definition of the **wait()** operation

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

- Definition of the **signal()** operation

```
signal(S) {  
    S++;  
}
```

Semaphore usage

- **Counting semaphore** – integer value can range over an unrestricted domain
- **Binary semaphore** – integer value can range only between 0 and 1
 - Same as a **mutex lock**
- Can solve various synchronization problems
- Consider P_1 and P_2 that require S_1 to happen before S_2
Create a semaphore “**synch**” initialized to 0
P1:
 S_1 ;
 signal(synch) ;
P2:
 wait(synch) ;
 S_2 ;
- Can implement a counting semaphore S as a binary semaphore

Semaphore Implementation

```
do
{
Wait (mutex);
//Critical section//
Signal(mutex);
//remainder section
}
While (true)
```

Semaphore Implementation

- With each semaphore there is an associated waiting queue
- Each entry in a waiting queue has two data items:
 - value (of type integer)
 - pointer to next record in the list
- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue
- ```
typedef struct{
 int value;

 struct process *list;

} semaphore;
```

# Semaphore implementation

```
wait(semaphore *S) {
 S->value--;
 if (S->value < 0) {
 add this process to S->list;
 block();
 }
}

signal(semaphore *S) {
 S->value++;
 if (S->value <= 0) {
 remove a process P from S->list;
 wakeup(P);
 }
}
```



# Implementation

- Semaphore operations now defined as

*wait(S):*

**S.value--;**

**if (S.value <= 0) {**

add this process to **S.L;**

**block;**

**}**

*signal(S):*

**S.value++;**

**if (S.value <= 0) {**

remove a process **P** from **S.L;**

**wakeup(P);**

**}**

# Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.
- Let  $S$  and  $Q$  be two semaphores initialized to 1

| $P_0$             | $P_1$             |
|-------------------|-------------------|
| <i>wait(S);</i>   | <i>wait(Q);</i>   |
| <i>wait(Q);</i>   | <i>wait(S);</i>   |
| $\vdots$          | $\vdots$          |
| <i>signal(S);</i> | <i>signal(Q);</i> |
| <i>signal(Q)</i>  | <i>signal(S);</i> |

- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.